



Project 2 : Data Classification

Datasets:

Yale Dataset

The dataset is taken from Original Yale Face Dataset (<http://vision.ucsd.edu/content/yale-face-database>). The Yale Face Database (size 6.4MB) contains 165 grayscale images in GIF format of 15 individuals. There are 11 images per subject, one per different facial expression or configuration: center-light, w/glasses, happy, left-light, w/no glasses, normal, right-light, sad, sleepy, surprised, and wink.

Some of the images are given below :



Handwriting Recognition Dataset(MNIST)

The MNIST database of handwritten digits, available from this page, has a training set of 60,000 examples, and a test set of 10,000 examples. It is a subset of a larger set available from NIST. The digits have been size-normalized and centered in a fixed-size image.

It is a good database for people who want to try learning techniques and pattern recognition methods on real-world data while spending minimal efforts on preprocessing and formatting.

Four files are available on this site:


train-images-idx3-ubyte.gz: training set images (9912422 bytes)

train-labels-idx1-ubyte.gz: training set labels (28881 bytes)

t10k-images-idx3-ubyte.gz: test set images (1648877 bytes)

t10k-labels-idx1-ubyte.gz: test set labels (4542 bytes)

Random Sampling of MNIST



Source Codes & their respective Results:

K-Nearest Neighbours (K = 3) :
(for Yale Dataset)

```
import os
import sys
import pdb
import glob
import math
import numpy
import random
import operator
from sets import Set
import scipy as scipy
from scipy.misc import *
from scipy import linalg

# This function loads the images, divides the data into training and test set
def getNeighbors(trainingSet, testInstance, trainingLabels, k):
    distances = []
    length = len(testInstance)-1

    for x in range(len(trainingSet)):
        #print(trainingSet[x])
        dist = EuclideanDistance(testInstance, trainingSet[x], length)
        distances.append((trainingSet[x], dist, trainingLabels[x]))
    distances.sort(key=operator.itemgetter(1))
    neighbors = []
    for x in range(k):
        #neighbors.append(distances[x][0])
        neighbors.append(distances[x][2])
    return neighbors

def getResponse(neighbors):
    classVotes = {}
    for x in range(len(neighbors)):
        response = neighbors[x]
        #print('response'+response)
        if response in classVotes:
            classVotes[response] += 1
        else:
            classVotes[response] = 1
    sortedVotes = sorted(classVotes.items(), key=operator.itemgetter(1), reverse=True)
    return sortedVotes[0][0]

def LoadImages(directory, split):
    # get a list of all the picture filenames
    gifs = glob.glob(directory + '/*.gif')
    # uncomment the below line when trying an unknown file
    #extraGif = glob.glob("/media/cosmos/Data/College Notes/M.Tech/Semester 4/Statistical Methods in AI/Project -
    Face Recognition/2.gif")
    classMap = {}
    testGIF = []
    allLabels = []
    testLabels = []
    trainingGIF = []
    trainingLabels = []
    for i in range(len(gifs)):
        if random.random() < split:
            trainingGIF.append(gifs[i])
            l = gifs[i].split("/")
```

```

        labelName = l[len(l)-1].split(".")[0][-2:]
        trainingLabels.append(labelName)
        allLabels.append(labelName)
    else:
        testGIF.append(gifs[i])
        l = gifs[i].split("/");
        labelName = l[len(l)-1].split(".")[0][-2:]
        testLabels.append(labelName)
        allLabels.append(labelName)
# uncomment the below 2 lines when trying an unknown file
#testGIF.append(extraGif[0])
#testLabels.append("un")
#allLabels.append("un")
trainingImgs = numpy.array([imread(i, True).flatten() for i in trainingGIF])
testImgs = numpy.array([imread(i, True).flatten() for i in testGIF])
# creating a list of class labels
allLabels = set(allLabels)
noOfClasses = len(allLabels)
sortedLabels = []
for i in allLabels:
    sortedLabels.append(i)
sortedLabels = sorted(sortedLabels)
# creating a mapping for confusion matrix
j = 0
for i in sortedLabels:
    classMap[i] = j
    j = j + 1
return trainingImgs, testImgs, trainingLabels, testLabels, noOfClasses, classMap

def EuclideanDistance(instance1, instance2, length):
    distance = 0
    for x in range(length):
        distance += pow((instance1[x] - instance2[x]), 2)
    return math.sqrt(distance)

def Mahanalobis(x, y):
    return scipy.spatial.distance.mahalanobis(x,y,np.linalg.inv(np.cov(x,y)))

# Run Principal Component Analysis on the input data.
# INPUT : data - an n x p matrix
# OUTPUT : e_faces -
#         weights -
#         mu -
def PCA(data):
    # mean
    mu = numpy.mean(data, 0)
    # mean adjust the data
    ma_data = data - mu
    # run SVD
    e_faces, sigma, v = linalg.svd(ma_data.transpose(), full_matrices=False)
    # compute weights for each image
    weights = numpy.dot(ma_data, e_faces)
    return e_faces, weights, mu

# This function calculates the weight of the test data
def InputWeight(testData, mu, e_faces):
    ma_data = testData - mu
    weights = numpy.dot(ma_data, e_faces)
    return weights

# Reconstruct an image using the given number of principal components.
def Reconstruct(imgIDx, e_faces, weights, mu, npcs):
    # dot weights with the eigenfaces and add to mean

```

```

recon = mu + numpy.dot(weights[imgIDX, 0:npcs], e_faces[:, 0:npcs].T)
return recon

# Saves the image in the given directory
def SaveImage(outDIR, subdir, imgID, imgDims, data):
    directory = outDIR + "/" + subdir
    if not os.path.exists(directory): os.makedirs(directory)
    imsave(directory + "/image_" + str(imgID) + ".jpg", data.reshape(imgDims))

# Prints the final results
def PrintResults(wrongPredictedClassCount, unknownLabels, accuracy, correctlyClassifiedDistances, maxDist,
confusionMatrix):
    print "Number of Wrongly Predicted Labels:",wrongPredictedClassCount
    print "Number of Unknown Labels:",unknownLabels
    print "Accuracy:",accuracy,"%"
    print "Max Distances among correctly classified:",correctlyClassifiedDistances[len(correctlyClassifiedDistances) -
1]
    print "Max Distances among all:",maxDist
    PrintConfusionMatrix(confusionMatrix)

# Prints the confusion matrix
def PrintConfusionMatrix(confusionMatrix):
    print "Confusion Matrix:"
    for i in range(0, len(confusionMatrix)):
        print confusionMatrix[i]

# Predicts the class of test data
def PredictLabelsFromTestData(testData, noOfClasses, mu, e_faces, trainingWeights, testLabels, trainingLabels,
classMap, thresholdDistance, noOfDimensions, k):
    correctlyClassifiedDistances = []
    confusionMatrix = [[0 for i in xrange(noOfClasses)] for i in xrange(noOfClasses)]
    wrongPredictedClassCount = 0
    unknownLabels = 0
    for i in range(len(testData)):
        testWeight=InputWeight(testData[i],mu,e_faces)
        distances = []
        for x in range(len(trainingWeights)):
            dist = EuclideanDistance(testWeight, trainingWeights[x], noOfDimensions)

            #dist = Mahanalobis(testWeight, trainingWeights[x])
            distances.append(dist)
        neighbors = getNeighbors(trainingWeights,testWeight, trainingLabels, k)
        result = getResponse(neighbors)
        predictedLabel = result
        actualLabel = testLabels[i]
        #predictedLabel = ""
        minDist = sys.maxint
        maxDist = -sys.maxint
        for j in range(len(distances)):
            if minDist > distances[j]:
                minDist = distances[j]
                predictedLabel = trainingLabels[j]
            if maxDist < distances[j]:
                maxDist = distances[j]
        confusionMatrix[classMap[actualLabel]][classMap[predictedLabel]] = confusionMatrix[classMap[actualLabel]]
[classMap[predictedLabel]] + 1
        #print "Actual class:",actualLabel
        #print "Predicted class:",predictedLabel
        #print "Min Dist:",minDist
        #print "-----"
        if minDist >= thresholdDistance:
            predictedLabel = "Unknown"
            unknownLabels = unknownLabels + 1

```

```

        elif predictedLabel != actualLabel:
            wrongPredictedClassCount = wrongPredictedClassCount + 1
        else:
            correctlyClassifiedDistances.append(minDist)
    # calculate accuracy
    accuracy = (1 - wrongPredictedClassCount / float(len(testData))) * 100
    correctlyClassifiedDistances.sort()
    PrintResults(wrongPredictedClassCount, unknownLabels, accuracy, correctlyClassifiedDistances, maxDist,
confusionMatrix)
    return accuracy

def main(arg):
    inDIR = "/home/cosmos/CSSstuff/SMAI/Project 2 - Classification/input/yalefaces"
    outDIR = "/home/cosmos/CSSstuff/SMAI/Project 2 - Classification/output/Yalefaces Output"
    k = int(arg[0])
    imgDims = (243, 320)
    split = 0.8
    noOfDimensions = 100
    thresholdDistance = 18000.0
    trainingData, testData, trainingLabels, testLabels, noOfClasses, classMap = LoadImages(inDIR, split)
    e_faces, trainingWeights, mu = PCA(trainingData)
    # save mean photo
    imsave(outDIR + "/mean.gif", mu.reshape(imgDims))
    # save each eigenface as an image
    for i in range(e_faces.shape[1]):
        SaveImage(outDIR, "eigenfaces", i, imgDims, e_faces[:,i])
    # reconstruct each face image using an increasing number of principal components
    reconstructed = []
    for p in range(trainingData.shape[0]):
        reconstructed.append(Reconstruct(p, e_faces, trainingWeights, mu, noOfDimensions))
        imgID = p
        SaveImage(outDIR, "reconstructed/" + str(p), imgID, imgDims, reconstructed[p])
    # Predicting Classes for test data
    accuracy = PredictLabelsFromTestData(testData, noOfClasses, mu, e_faces, trainingWeights, testLabels,
trainingLabels, classMap, thresholdDistance, noOfDimensions, k)
    return accuracy

total = 0.0
for i in range(0, 5):
    total = total + main(sys.argv[1:])
    print "_____ "
print "Mean Accuracy:", total / 5, "%"

```

Results:

Number of Wrongly Predicted Labels: 5
 Number of Unknown Labels: 0
 Accuracy: 84.8484848485 %
 Max Distances among correctly classified: 13029.804586
 Max Distances among all: 44376.4745935
 Confusion Matrix:
 [2, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
 [0, 4, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0]
 [0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
 [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
 [0, 0, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
 [0, 0, 0, 0, 0, 2, 0, 1, 0, 0, 0, 0, 0, 0, 0]
 [0, 0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 1]
 [0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0]
 [0, 0, 0, 0, 0, 0, 0, 0, 4, 0, 0, 0, 0, 0, 0]
 [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
 [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0]

[0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 3, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2]

Number of Wrongly Predicted Labels: 3
Number of Unknown Labels: 0
Accuracy: 90.3225806452 %
Max Distances among correctly classified: 14559.4523311
Max Distances among all: 37696.4649247

Confusion Matrix:

[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 3, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 4, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 2]
[0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 3, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 3]

Number of Wrongly Predicted Labels: 4
Number of Unknown Labels: 0
Accuracy: 88.2352941176 %
Max Distances among correctly classified: 17670.4641544
Max Distances among all: 37367.8789651

Confusion Matrix:

[2, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 3, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 4, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 1, 2, 0, 0, 0, 0, 0, 1]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 3, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 4, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 3]

Number of Wrongly Predicted Labels: 5
Number of Unknown Labels: 0
Accuracy: 83.8709677419 %
Max Distances among correctly classified: 10415.4792513
Max Distances among all: 37176.1344234

Confusion Matrix:

[4, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[1, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0]

```
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 4, 0, 0, 0]
[0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 3, 0, 0]
[0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 2, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 3, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 3]
```

Number of Wrongly Predicted Labels: 6

Number of Unknown Labels: 0

Accuracy: 83.3333333333 %

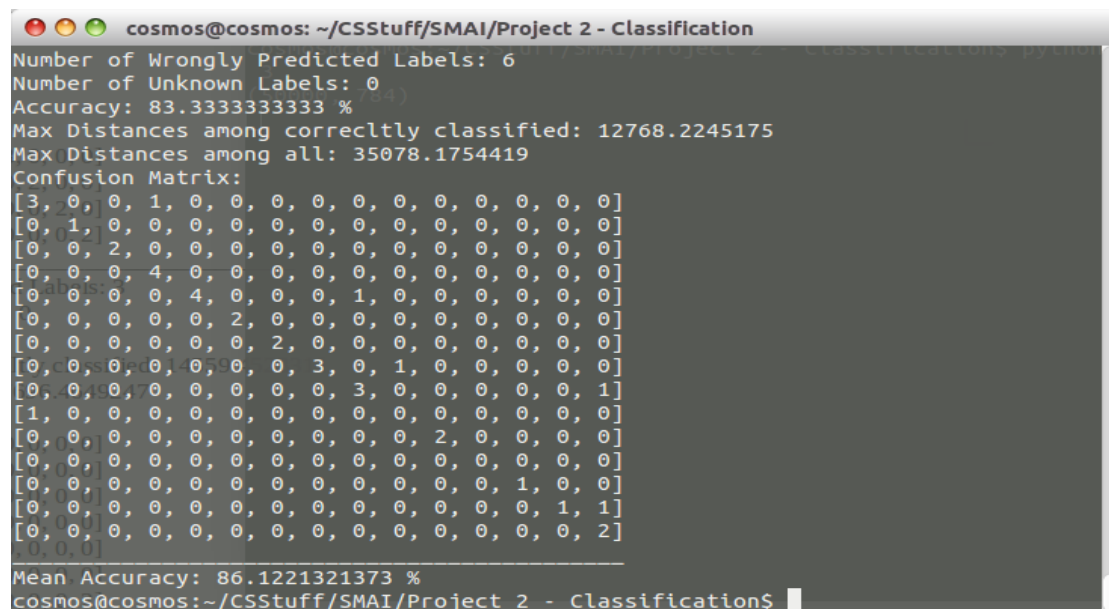
Max Distances among correctly classified: 12768.2245175

Max Distances among all: 35078.1754419

Confusion Matrix:

```
[3, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 4, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 4, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 3, 0, 1, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 3, 0, 0, 0, 0, 0, 1]
[1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2]
```

Mean Accuracy: 86.1221321373 %

A screenshot of a terminal window titled "cosmos@cosmos: ~/CSStuff/SMAI/Project 2 - Classification". The terminal displays the same classification results as the text above, including the list of 15 predicted labels, the number of wrongly predicted labels (6), unknown labels (0), accuracy (83.3333333333 %), max distances, and a 15x15 confusion matrix. At the bottom, it shows the mean accuracy (86.1221321373 %) and the prompt "cosmos@cosmos:~/CSStuff/SMAI/Project 2 - Classification\$".

```
cosmos@cosmos: ~/CSStuff/SMAI/Project 2 - Classification
Number of Wrongly Predicted Labels: 6
Number of Unknown Labels: 0
Accuracy: 83.3333333333 %
Max Distances among correctly classified: 12768.2245175
Max Distances among all: 35078.1754419
Confusion Matrix:
[3, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 4, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 4, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 3, 0, 1, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 3, 0, 0, 0, 0, 0, 1]
[1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2]
Mean Accuracy: 86.1221321373 %
cosmos@cosmos:~/CSStuff/SMAI/Project 2 - Classification$
```


For MNIST Dataset:

```
import mnist_loader
import os
import sys
import pdb
import glob
import math
import numpy
import random
import operator
from sets import Set
import scipy as scipy
from scipy.misc import *
from scipy import linalg

# Prints the confusion matrix
def getNeighbors(trainingSet, testInstance, trainingLabels, k):
    distances = []
    length = len(testInstance)-1

    for x in range(len(trainingSet)):
        #print(trainingSet[x])
        dist = EuclideanDistance(testInstance, trainingSet[x], length)
        distances.append((trainingSet[x], dist, trainingLabels[x]))
    distances.sort(key=operator.itemgetter(1))
    neighbors = []
    for x in range(k):
        neighbors.append(distances[x][2])
    return neighbors

def getResponse(neighbors):
    classVotes = {}
    for x in range(len(neighbors)):
        response = neighbors[x]
        #print('response'+response)
        if response in classVotes:
            classVotes[response] += 1
        else:
            classVotes[response] = 1
    sortedVotes = sorted(classVotes.items(), key=operator.itemgetter(1), reverse=True)
    return sortedVotes[0][0]

def CreateConfusionMatrix(predictions, testSet):
    noOfClasses = 10
    confusionMatrix = [[0 for i in xrange(noOfClasses)] for i in xrange(noOfClasses)]
    s=set()
    for i in testSet:
        s.add(i)
    l=list(s)
    d={}
    for x in range(len(l)):
        d[l[x]]=x

    for x in range(len(testSet)):
        if predictions[x]==testSet[x]:
            confusionMatrix[d[testSet[x]]][d[testSet[x]]] = confusionMatrix[d[testSet[x]]][d[testSet[x]]] + 1
        else:
            confusionMatrix[d[testSet[x]]][d[predictions[x]]] = confusionMatrix[d[testSet[x]]][d[predictions[x]]] + 1
    return confusionMatrix

def CalculatePrecisionAndRecall(confusionMatrix, noOfClasses, noOfTestSamples):
    totalRecall = 0.0
```

```

totalPrecision = 0.0
totalSpecificity = 0.0
recall = []
precision = []
specificity = []
#print "Precisions for Different Classes"
#print "_____ "
for i in range(0, len(confusionMatrix[0])):
    classPrecision = 0.0
    for j in range(0, len(confusionMatrix)):
        classPrecision = classPrecision + confusionMatrix[j][i]
    if classPrecision != 0.0:
        classPrecision = (confusionMatrix[i][i] / float(classPrecision)) * 100
    else:
        classPrecision = 0.0
    precision.append(classPrecision)
    #print "Class Precision for class", i + 1, ":", classPrecision
    totalPrecision = totalPrecision + classPrecision
#print "Recalls for Different Classes"
#print "_____ "
for i in range(0, len(confusionMatrix)):
    classRecall = 0.0
    for j in range(0, len(confusionMatrix[i])):
        classRecall = classRecall + confusionMatrix[i][j]
    if classRecall != 0.0:
        classRecall = (confusionMatrix[i][i] / float(classRecall)) * 100
    else:
        classRecall = 0.0
    recall.append(classRecall)
    #print "Class Recall for class", i + 1, ":", classRecall
    totalRecall = totalRecall + classRecall
for i in range(0, len(confusionMatrix[0])):
    numerator = noOfTestSamples - confusionMatrix[i][i]
    denominator = numerator
    for j in range(0, len(confusionMatrix)):
        if i != j:
            denominator = denominator + confusionMatrix[j][i]
    classSpecificity = (numerator / float(denominator)) * 100
    totalSpecificity = totalSpecificity + classSpecificity
    specificity.append(classSpecificity)

avgRecall = (totalRecall / float(noOfClasses))
avgPrecision = (totalPrecision / float(noOfClasses))
avgSpecificity = (totalSpecificity / float(noOfClasses))
return avgPrecision, avgRecall, avgSpecificity, precision, recall, specificity

```

```

def PrintResults(confusionMatrix, avgPrecision, avgRecall, avgSpecificity, precision, recall, specificity):
    print "Confusion Matrix:"
    for i in range(0, len(confusionMatrix)):
        print confusionMatrix[i]
    for i in range(0, len(precision)):
        print "Class", i + 1
        print "-----"
        print "Precision :", precision[i]
        print "Recall :", recall[i]
        print "Specificity :", specificity[i]
        print "\n"
    print "-----"
    print "Average Recall:", avgRecall
    print "Average Precision:", avgPrecision
    print "Average Specificity:", avgSpecificity

```

```

def EuclideanDistance(instance1, instance2, length):

```

```

        distance = 0
        for x in range(length):
            distance += pow((instance1[x] - instance2[x]), 2)
        return math.sqrt(distance)

def GetAccuracy(testLabels, predictions):
    correct = 0

    for x in range(len(testLabels)):
        if testLabels[x] == predictions[x]:
            correct += 1

    return (correct/float(len(testLabels))) * 100.0

def main(k):
    training_data, validation_data, test_data = mnist_loader.load_data()
    k = int(k[0])
    predictions=[]
    for i in range( len(test_data[0])):
        neighbors=getNeighbors(training_data[0],test_data[0][i],training_data[1],k)
        result=getResponse(neighbors)
        predictions.append(result)
    confusionMatrix = CreateConfusionMatrix(predictions,test_data[1])
    avgPrecision, avgRecall, avgSpecificity, precision, recall, specificity =
CalculatePrecisionAndRecall(confusionMatrix, 10, len(test_data[1]))
    PrintResults(confusionMatrix, avgPrecision, avgRecall, avgSpecificity, precision, recall, specificity)
    accuracy = GetAccuracy(test_data[1], predictions)
    print "Accuracy :", accuracy

main(sys.argv[1:])

```

Results:

Confusion Matrix:

```

[967, 0, 1, 0, 0, 5, 4, 1, 2, 0]
[0, 1120, 2, 3, 0, 1, 3, 1, 5, 0]
[9, 1, 962, 7, 10, 1, 13, 11, 16, 2]
[1, 1, 14, 950, 1, 17, 1, 10, 11, 4]
[1, 1, 7, 0, 937, 0, 7, 2, 2, 25]
[7, 4, 5, 33, 7, 808, 11, 2, 10, 5]
[10, 3, 4, 1, 5, 10, 924, 0, 1, 0]
[2, 13, 22, 5, 7, 1, 0, 954, 4, 20]
[4, 6, 6, 14, 8, 24, 10, 8, 891, 3]
[10, 6, 0, 12, 33, 5, 1, 14, 6, 922]

```

Class 1

```

Precision : 95.6478733927
Recall : 98.6734693878
Specificity : 99.5152583453

```

Class 2

```

Precision : 96.9696969697
Recall : 98.6784140969
Specificity : 99.6074032529

```

Class 3

```

Precision : 94.03714565
Recall : 93.2170542636

```

Specificity : 99.329596659

Class 4

Precision : 92.6829268293
Recall : 94.0594059406
Specificity : 99.1780821918

Class 5

Precision : 92.9563492063
Recall : 95.4175152749
Specificity : 99.2226844756

Class 6

Precision : 92.6605504587
Recall : 90.5829596413
Specificity : 99.3085566119

Class 7

Precision : 94.8665297741
Recall : 96.4509394572
Specificity : 99.4521148367

Class 8

Precision : 95.1146560319
Recall : 92.8015564202
Specificity : 99.4612424409

Class 9

Precision : 93.9873417722
Recall : 91.4784394251
Specificity : 99.3781365918

Class 10

Precision : 93.9857288481
Recall : 91.3776015857
Specificity : 99.3542738317

Average Recall: 94.2737355493
Average Precision: 94.2908798933
Average Specificity: 99.3807349238
Accuracy : 82.15

SVM:

For Yale Dataset:

1. Source Code

```
import sys
import glob
import numpy
import random
from scipy.misc import *
from scipy import linalg
from subprocess import call

def SplitData(directory):
    split = 0.8
    gifs = glob.glob(directory + '/*.gif')
    classMap = {}
    testGIF = []
    allLabels = []
    testLabels = []
    trainingGIF = []
    trainingLabels = []
    for i in range(len(gifs)):
        if random.random() < split:
            trainingGIF.append(gifs[i])
            l = gifs[i].split("/")
            labelName = l[len(l)-1].split(".")[0][-2:]
            trainingLabels.append(labelName)
            allLabels.append(labelName)
        else:
            testGIF.append(gifs[i])
            l = gifs[i].split("/")
            labelName = l[len(l)-1].split(".")[0][-2:]
            testLabels.append(labelName)
            allLabels.append(labelName)
    trainingImgs = numpy.array([imread(i, True).flatten() for i in trainingGIF])
    testImgs = numpy.array([imread(i, True).flatten() for i in testGIF])
    allLabels = set(allLabels)
    noOfClasses = len(allLabels)
    sortedLabels = []
    for i in allLabels:
        sortedLabels.append(i)
    sortedLabels = sorted(sortedLabels)
    # creating a mapping for confusion matrix
    j = 0
    for i in sortedLabels:
        if i[0] == '0':
            i = i[1 : ]
        classMap[i] = j
        j = j + 1
    return trainingImgs, testImgs, trainingLabels, testLabels, noOfClasses, classMap

def PCA(data):
    # mean
    mu = numpy.mean(data, 0)
    # mean adjust the data
    ma_data = data - mu
    # run SVD
    e_faces, sigma, v = linalg.svd(ma_data.transpose(), full_matrices=False)
    # compute weights for each image
    weights = numpy.dot(ma_data, e_faces)
    return e_faces, weights, mu
```

```

def WriteIntoFile(data, labels, directory, fileType):
    if fileType == "train":
        fd = open(directory + "/training.txt", 'w+')
    else:
        fd = open(directory + "/test.txt", 'w+')
    for i in range(0, len(data)):
        line = labels[i]
        for j in range(0, len(data[i])):
            line = line + " " + str(j + 1) + ":" + str(data[i][j])
        fd.write(line + "\n")
    fd.close()

def TrainUsingSVM(directory):
    filePath = directory + "/training.txt"
    command = directory + "/svm-train"
    call([command, "-t", "0", filePath])

def InputWeight(testData, mu, e_faces):
    ma_data = testData - mu
    weights = numpy.dot(ma_data, e_faces)
    return weights

def PredictLabels(directory):
    command = directory + "output/svm-predict"
    testFilePath = directory + "output/test.txt"
    modelPath = directory + "/training.txt.model"
    fd = open(directory + "/output/result", 'w+')
    resultFilePath = directory + "output/result"
    call([command, testFilePath, modelPath, resultFilePath])
    fd.close()

def CalculateAccuracy(directory, actualLabels):
    noOfCorrectlyClassifiedSamples = 0
    fd = open(directory + "/result", "r")
    predictedLabels = []
    for line in fd:
        predictedLabels.append(line)
    for i in range(0, len(predictedLabels)):
        predictedLabels[i] = predictedLabels[i][0 : len(predictedLabels[i]) - 1]
    for i in range(0, len(actualLabels)):
        if actualLabels[i][0] == '0':
            actualLabels[i] = actualLabels[i][1 : ]
    for i in range(0, len(actualLabels)):
        if actualLabels[i] == predictedLabels[i]:
            noOfCorrectlyClassifiedSamples = noOfCorrectlyClassifiedSamples + 1
    accuracy = (noOfCorrectlyClassifiedSamples / float(len(actualLabels))) * 100
    return actualLabels, predictedLabels, accuracy

def CalculateConfusionMatrix(actualLabels, predictedLabels, classMap, noOfClasses):
    confusionMatrix = [[0 for i in xrange(noOfClasses)] for i in xrange(noOfClasses)]
    for i in range(0, len(actualLabels)):
        confusionMatrix[classMap[actualLabels[i]]][classMap[predictedLabels[i]]] =
confusionMatrix[classMap[actualLabels[i]]][classMap[predictedLabels[i]]] + 1
    return confusionMatrix

def CalculatePrecisionAndRecall(confusionMatrix, noOfClasses, noOfTestSamples):
    totalRecall = 0.0
    totalPrecision = 0.0
    totalSpecificity = 0.0
    recall = []
    precision = []
    specificity = []

```

```

#print "Precisions for Different Classes"
#print "_____ "
for i in range(0, len(confusionMatrix[0])):
    classPrecision = 0.0
    for j in range(0, len(confusionMatrix)):
        classPrecision = classPrecision + confusionMatrix[j][i]
    if classPrecision != 0.0:
        classPrecision = (confusionMatrix[i][i] / float(classPrecision)) * 100
    else:
        classPrecision = 0.0
    precision.append(classPrecision)
    #print "Class Precision for class", i + 1, ":", classPrecision
    totalPrecision = totalPrecision + classPrecision
#print "Recalls for Different Classes"
#print "_____ "
for i in range(0, len(confusionMatrix)):
    classRecall = 0.0
    for j in range(0, len(confusionMatrix[i])):
        classRecall = classRecall + confusionMatrix[i][j]
    if classRecall != 0.0:
        classRecall = (confusionMatrix[i][i] / float(classRecall)) * 100
    else:
        classRecall = 0.0
    recall.append(classRecall)
    #print "Class Recall for class", i + 1, ":", classRecall
    totalRecall = totalRecall + classRecall
for i in range(0, len(confusionMatrix[0])):
    numerator = noOfTestSamples - confusionMatrix[i][i]
    denominator = numerator
    for j in range(0, len(confusionMatrix)):
        if i != j:
            denominator = denominator + confusionMatrix[j][i]
    classSpecificity = (numerator / float(denominator)) * 100
    totalSpecificity = totalSpecificity + classSpecificity
    specificity.append(classSpecificity)

avgRecall = (totalRecall / float(noOfClasses))
avgPrecision = (totalPrecision / float(noOfClasses))
avgSpecificity = (totalSpecificity / float(noOfClasses))
return avgPrecision, avgRecall, avgSpecificity, precision, recall, specificity

```

```

def PrintResults(confusionMatrix, avgPrecision, avgRecall, avgSpecificity, precision, recall, specificity, accuracy,
iterationNo):

```

```

    print "Confusion Matrix:"
    for i in range(0, len(confusionMatrix)):
        print confusionMatrix[i]
    for i in range(0, len(precision)):
        print "Class", i + 1
        print "-----"
        print "Precision :", precision[i]
        print "Recall :", recall[i]
        print "Specificity :", specificity[i]
        print "\n"
    print "-----"
    print "Average Recall:", avgRecall
    print "Average Precision:", avgPrecision
    print "Average Specificity:", avgSpecificity
    print "Accuracy for iteration number", iterationNo + 1, ":", accuracy

```

```

def main(argv, iterationNo):
    directory = argv[0]
    inDIR = directory + "input/yalefaces"
    outDIR = directory + "output"

```

```

trainingData, testData, trainingLabels, testLabels, noOfClasses, classMap = SplitData(inDIR)
e_faces, trainingWeights, mu = PCA(trainingData)
WriteIntoFile(trainingWeights, trainingLabels, outDIR, "train")
TrainUsingSVM(outDIR)
testWeights = InputWeight(testData, mu, e_faces)
WriteIntoFile(testWeights, testLabels, outDIR, "test")
PredictLabels(directory)
actualLabels, predictedLabels, accuracy = CalculateAccuracy(outDIR, testLabels)
confusionMatrix = CalculateConfusionMatrix(actualLabels, predictedLabels, classMap, noOfClasses)
avgPrecision, avgRecall, avgSpecificity, precision, recall, specificity =
CalculatePrecisionAndRecall(confusionMatrix, noOfClasses, len(testData))
PrintResults(confusionMatrix, avgPrecision, avgRecall, avgSpecificity, precision, recall, specificity, accuracy,
iterationNo)
return accuracy

total = 0.0
for i in range(0, 5):
    accuracy = main(sys.argv[1:], i)
    total = total + accuracy
    print "_____ "
print "Average Accuracy : ",total / 5.0,"%"
```

2. Results

Accuracy = 92.1053% (35/38) (classification)

Confusion Matrix:

```

[5, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 3, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0]
[0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 5, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 5, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 2, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2]
```

Class 1

Precision : 100.0

Recall : 100.0

Specificity : 100.0

Class 2

Precision : 100.0

Recall : 100.0

Specificity : 100.0

Class 3

Precision : 66.6666666667

Recall : 100.0

Specificity : 97.2972972973

Class 4

Precision : 100.0

Recall : 100.0

Specificity : 100.0

Class 5

Precision : 100.0

Recall : 33.3333333333

Specificity : 100.0

Class 6

Precision : 100.0

Recall : 100.0

Specificity : 100.0

Class 7

Precision : 50.0

Recall : 100.0

Specificity : 97.3684210526

Class 8

Precision : 100.0

Recall : 100.0

Specificity : 100.0

Class 9

Precision : 100.0

Recall : 100.0

Specificity : 100.0

Class 10

Precision : 100.0

Recall : 100.0

Specificity : 100.0

Class 11

Precision : 100.0

Recall : 100.0

Specificity : 100.0

Class 12

Precision : 50.0

Recall : 100.0

Specificity : 97.3684210526

Class 13

Precision : 100.0

Recall : 66.6666666667

Specificity : 100.0

Class 14

Precision : 100.0

Recall : 100.0

Specificity : 100.0

Class 15

Precision : 100.0

Recall : 100.0

Specificity : 100.0

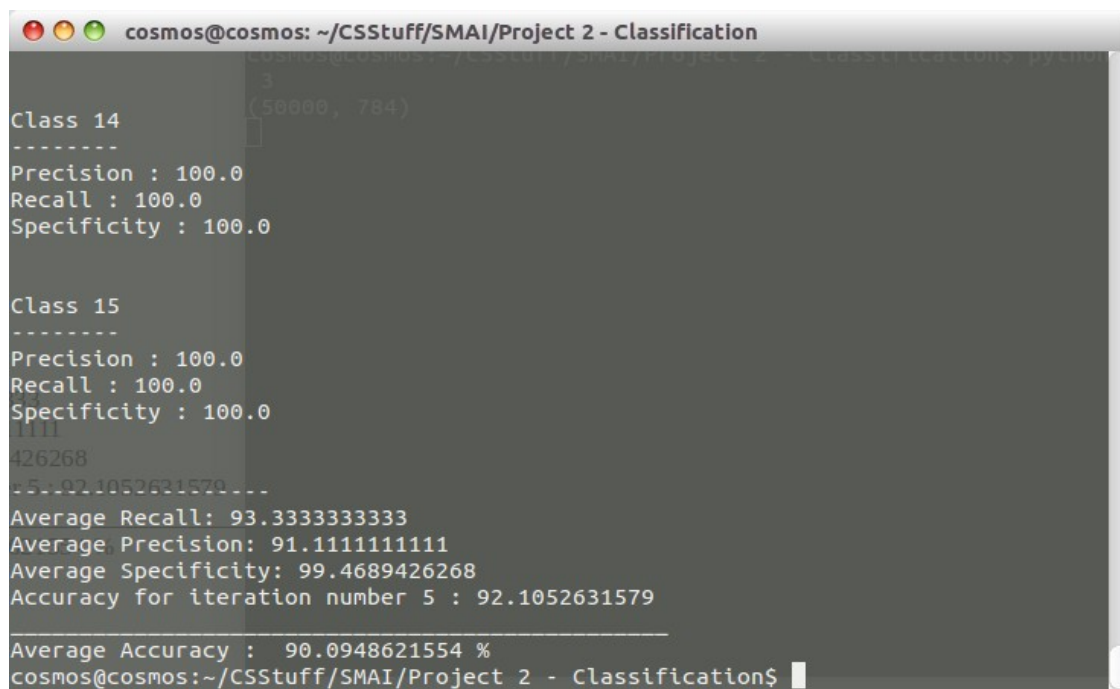
Average Recall: 93.3333333333

Average Precision: 91.1111111111

Average Specificity: 99.4689426268

Accuracy for iteration number 5 : 92.1052631579

Average Accuracy : 90.0948621554 %



```
cosmos@cosmos: ~/CSStuff/SMAl/Project 2 - Classification
3
(50000, 784)
Class 14
-----
Precision : 100.0
Recall : 100.0
Specificity : 100.0

Class 15
-----
Precision : 100.0
Recall : 100.0
Specificity : 100.0
1111
426268
92.1052631579
Average Recall: 93.3333333333
Average Precision: 91.1111111111
Average Specificity: 99.4689426268
Accuracy for iteration number 5 : 92.1052631579

Average Accuracy : 90.0948621554 %
cosmos@cosmos:~/CSStuff/SMAl/Project 2 - Classification$
```

For MNIST Dataset:

1. Source Code

```
"""  
mnist_svm  
~~~~~
```

A classifier program for recognizing handwritten digits from the MNIST data set, using an SVM classifier. """

```
#### Libraries  
# My libraries  
import mnist_loader  
  
# Third-party libraries  
from sklearn import svm  
  
def CreateConfusionMatrix(predictions, testSet):  
    noOfClasses = 10  
    confusionMatrix = [[0 for i in xrange(noOfClasses)] for i in xrange(noOfClasses)]  
    s=set()  
    for i in testSet:  
        s.add(i)  
    l=list(s)  
    d={}  
    for x in range(len(l)):  
        d[l[x]]=x  
  
    for x in range(len(testSet)):  
        if predictions[x]==testSet[x]:  
            confusionMatrix[d[testSet[x]]][d[testSet[x]]] = confusionMatrix[d[testSet[x]]][d[testSet[x]]] + 1  
        else:  
            confusionMatrix[d[testSet[x]]][d[predictions[x]]] = confusionMatrix[d[testSet[x]]][d[predictions[x]]] + 1  
    return confusionMatrix  
  
def CalculatePrecisionAndRecall(confusionMatrix, noOfClasses, noOfTestSamples):  
    totalRecall = 0.0  
    totalPrecision = 0.0  
    totalSpecificity = 0.0  
    recall = []  
    precision = []  
    specificity = []  
    #print "Precisions for Different Classes"  
    #print " _____ "  
    for i in range(0, len(confusionMatrix[0])):  
        classPrecision = 0.0  
        for j in range(0, len(confusionMatrix)):  
            classPrecision = classPrecision + confusionMatrix[j][i]  
        if classPrecision != 0.0:  
            classPrecision = (confusionMatrix[i][i] / float(classPrecision)) * 100  
        else:  
            classPrecision = 0.0  
        precision.append(classPrecision)  
        #print "Class Precision for class", i + 1, ":", classPrecision  
        totalPrecision = totalPrecision + classPrecision  
    #print "Recalls for Different Classes"  
    #print " _____ "  
    for i in range(0, len(confusionMatrix)):  
        classRecall = 0.0  
        for j in range(0, len(confusionMatrix[i])):  
            classRecall = classRecall + confusionMatrix[i][j]
```

```

        if classRecall != 0.0:
            classRecall = (confusionMatrix[i][i] / float(classRecall)) * 100
        else:
            classRecall = 0.0
        recall.append(classRecall)
        #print "Class Recall for class", i + 1, ":", classRecall
        totalRecall = totalRecall + classRecall
    for i in range(0, len(confusionMatrix[0])):
        numerator = noOfTestSamples - confusionMatrix[i][i]
        denominator = numerator
        for j in range(0, len(confusionMatrix)):
            if i != j:
                denominator = denominator + confusionMatrix[j][i]
        classSpecificity = (numerator / float(denominator)) * 100
        totalSpecificity = totalSpecificity + classSpecificity
    specificity.append(classSpecificity)

avgRecall = (totalRecall / float(noOfClasses))
avgPrecision = (totalPrecision / float(noOfClasses))
avgSpecificity = (totalSpecificity / float(noOfClasses))
return avgPrecision, avgRecall, avgSpecificity, precision, recall, specificity

def PrintResults(confusionMatrix, avgPrecision, avgRecall, avgSpecificity, precision, recall, specificity):
    print "Confusion Matrix:"
    for i in range(0, len(confusionMatrix)):
        print confusionMatrix[i]
    for i in range(0, len(precision)):
        print "Class", i + 1
        print "-----"
        print "Precision :", precision[i]
        print "Recall :", recall[i]
        print "Specificity :", specificity[i]
        print "\n"
    print "-----"
    print "Average Recall:", avgRecall
    print "Average Precision:", avgPrecision
    print "Average Specificity:", avgSpecificity

def svm_baseline():
    training_data, validation_data, test_data = mnist_loader.load_data()

    # train
    clf = svm.SVC()
    clf.fit(training_data[0], training_data[1])
    # test
    predictions = [int(a) for a in clf.predict(test_data[0])]
    num_correct = sum(int(a == y) for a, y in zip(predictions, test_data[1]))
    confusionMatrix = CreateConfusionMatrix(predictions, test_data[1])
    avgPrecision, avgRecall, avgSpecificity, precision, recall, specificity =
    CalculatePrecisionAndRecall(confusionMatrix, 10, len(test_data[1]))
    PrintResults(confusionMatrix, avgPrecision, avgRecall, avgSpecificity, precision, recall, specificity)
    print "Baseline classifier using an SVM."
    print "%s of %s values correct." % (num_correct, len(test_data[1]))

if __name__ == "__main__":
    svm_baseline()

```

2. Results

Confusion Matrix:

```
[967, 0, 1, 0, 0, 5, 4, 1, 2, 0]
[0, 1120, 2, 3, 0, 1, 3, 1, 5, 0]
[9, 1, 962, 7, 10, 1, 13, 11, 16, 2]
[1, 1, 14, 950, 1, 17, 1, 10, 11, 4]
[1, 1, 7, 0, 937, 0, 7, 2, 2, 25]
[7, 4, 5, 33, 7, 808, 11, 2, 10, 5]
[10, 3, 4, 1, 5, 10, 924, 0, 1, 0]
[2, 13, 22, 5, 7, 1, 0, 954, 4, 20]
[4, 6, 6, 14, 8, 24, 10, 8, 891, 3]
[10, 6, 0, 12, 33, 5, 1, 14, 6, 922]
```

Class 1

Precision : 95.6478733927
Recall : 98.6734693878
Specificity : 99.5152583453

Class 2

Precision : 96.9696969697
Recall : 98.6784140969
Specificity : 99.6074032529

Class 3

Precision : 94.03714565
Recall : 93.2170542636
Specificity : 99.329596659

Class 4

Precision : 92.6829268293
Recall : 94.0594059406
Specificity : 99.1780821918

Class 5

Precision : 92.9563492063
Recall : 95.4175152749
Specificity : 99.2226844756

Class 6

Precision : 92.6605504587
Recall : 90.5829596413
Specificity : 99.3085566119

Class 7

Precision : 94.8665297741
Recall : 96.4509394572
Specificity : 99.4521148367

Class 8

Precision : 95.1146560319
Recall : 92.8015564202
Specificity : 99.4612424409

Class 9

Precision : 93.9873417722
Recall : 91.4784394251
Specificity : 99.3781365918

Class 10

Precision : 93.9857288481
Recall : 91.3776015857
Specificity : 99.3542738317

Average Recall: 94.2737355493
Average Precision: 94.2908798933
Average Specificity: 99.3807349238
Baseline classifier using an SVM.
9435 of 10000 values correct.

```
Recall : 90.5829596413
Specificity : 99.3085566119

def PrintResults(confusionMatrix, avgPrecision, avgRecall, avgSpecificity, precision, recall, specificity):
    print("Confusion Matrix:")
    for i in range(9, len(confusionMatrix)):
        print("Precision : %4.8665297741" % precision[i])
        print("Recall : %4.4509394572" % recall[i])
        print("Specificity : %4.4521148367" % specificity[i])
    print("Precision : %4.8665297741" % precision[i])
    print("Recall : %4.4509394572" % recall[i])
    print("Specificity : %4.4521148367" % specificity[i])

Class 8
Precision : 95.1146560319
Recall : 92.8015564202
Specificity : 99.4612424409

Class 9
Precision : 93.9873417722
Recall : 91.4784394251
Specificity : 99.3781365918

Class 10
Precision : 93.9857288481
Recall : 91.3776015857
Specificity : 99.3542738317

Average Recall: 94.2737355493
Average Precision: 94.2908798933
Average Specificity: 99.3807349238
Baseline classifier using an SVM.
9435 of 10000 values correct.
```

Multilayer Feedforward Neural Network

For Yale Dataset:

1. Source Code

`eigen_face_loader`

```
# -*- coding: utf-8 -*-  
"""
```

Created on Thu Apr 9 22:59:21 2015

```
@author: pramod  
"""
```

```
import numpy as numpy  
import random  
import glob  
from scipy.misc import *
```

```
def LoadImages(directory, split):  
    # get a list of all the picture filenames  
    gifs = glob.glob(directory + '/*.gif')  
    # uncomment the below line when trying an unknown file  
    #extraGif = glob.glob("/media/cosmos/Data/College Notes/M.Tech/Semester 4/Statistical Methods in AI/Project -  
Face Recognition/2.gif")  
    classMap = {}  
    testGIF = []  
    allLabels = []  
    testLabels = []  
    trainingGIF = []  
    trainingLabels = []  
    for i in range(len(gifs)):  
        if random.random() < split:  
            trainingGIF.append(gifs[i])  
            l = gifs[i].split("/")  
            labelName = l[len(l)-1].split(".")[0][-2:]  
            trainingLabels.append(labelName)  
            allLabels.append(labelName)  
        else:  
            testGIF.append(gifs[i])  
            l = gifs[i].split("/")  
            labelName = l[len(l)-1].split(".")[0][-2:]  
            testLabels.append(labelName)  
            allLabels.append(labelName)  
    # uncomment the below 2 lines when trying an unknown file  
    #testGIF.append(extraGif[0])  
    #testLabels.append("un")  
    #allLabels.append("un")  
    trainingImgs = numpy.array([imread(i, True).flatten() for i in trainingGIF])  
    testImgs = numpy.array([imread(i, True).flatten() for i in testGIF])  
    # creating a list of class labels  
    allLabels = set(allLabels)  
    noOfClasses = len(allLabels)  
    sortedLabels = []  
    for i in allLabels:  
        sortedLabels.append(i)  
    sortedLabels = sorted(sortedLabels)  
    # creating a mapping for confusion matrix  
    j = 0  
    for i in sortedLabels:  
        classMap[i] = j  
        j = j + 1  
    return trainingImgs, testImgs, trainingLabels, testLabels, noOfClasses, classMap  
# mean
```

```

mu = numpy.mean(data, 0)
# mean adjust the data
ma_data = data - mu
# run SVD
e_faces, sigma, v = numpy.linalg.svd(ma_data.transpose(), full_matrices=False)
# compute weights for each image
weights = numpy.dot(ma_data, e_faces)
return e_faces, weights, mu
def PCA(data):
    # mean
    mu = numpy.mean(data, 0)
    # mean adjust the data
    ma_data = data - mu
    # run SVD
    e_faces, sigma, v = numpy.linalg.svd(ma_data.transpose(), full_matrices=False)
    # compute weights for each image
    weights = numpy.dot(ma_data, e_faces)
    return e_faces, weights, mu
def InputWeight(testData, mu, e_faces):
    ma_data = testData - mu
    weights = numpy.dot(ma_data, e_faces)
    return weights
def load_data():
    inDIR = "/home/cosmos/CSSstuff/SMAI/Project 2 - Classification/input/yalefaces"
    outDIR = "/home/cosmos/CSSstuff/SMAI/Project 2 - Classification/output"
    imgDims = (243, 320)

    split = 0.8
    trainingData, testData, trainingLabels, testLabels, noOfClasses, classMap = LoadImages(inDIR, split)
    e_faces, trainingWeights, mu = PCA(trainingData)
    #print trainingWeights.shape[0]
    #print trainingWeights.shape
    #print mu.shape
    #print e_faces.shape
    """
    tr_d, va_d, te_d = load_data()
    training_inputs = [np.reshape(x, (784, 1)) for x in tr_d[0]]
    print tr_d[0].shape
    training_results = [vectorized_result(y) for y in tr_d[1]]
    training_data = zip(training_inputs, training_results)
    validation_inputs = [np.reshape(x, (784, 1)) for x in va_d[0]]
    validation_data = zip(validation_inputs, va_d[1])
    test_inputs = [np.reshape(x, (784, 1)) for x in te_d[0]]
    test_data = zip(test_inputs, te_d[1])
    return (training_data, validation_data, test_data)
    """

    testWeights=numpy.zeros((len(testData),trainingWeights.shape[1]))
    #print testWeights.shape
    #print trainingWeights.shape
    #print len(testWeights)
    for i in range(len(testWeights)):
        #testWeight =InputWeight(testData[i],mu,e_faces)
        testWeights[i]=InputWeight(testData[i],mu,e_faces)

    #formatting weights
    training_inputs = [numpy.reshape(x, (len(trainingWeights), 1)) for x in trainingWeights]
    test_inputs= [numpy.reshape(x,(len(trainingWeights) , 1)) for x in testWeights]
    #Convert training labels to vectors
    trainingLabels=numpy.asarray(trainingLabels)
    trainingLabels=trainingLabels.astype(numpy.float)
    testLabels=numpy.asarray(testLabels)
    testLabels=testLabels.astype(numpy.float)
    #decrementing subject labels

```



```

for i in range(len(trainingLabels)):
    trainingLabels[i]=trainingLabels[i]-1
for i in range(len(testLabels)):
    testLabels[i]=testLabels[i]-1
#print testLabels
#print vectorized_result(trainingLabels[0])
#print trainingWeights.shape
training_results = [vectorized_result(y) for y in trainingLabels]
#test_results= [vectorized_result(y) for y in testLabels]

tr_d=zip(training_inputs, training_results)
te_d=zip(test_inputs,testLabels)
#print te_d
return (tr_d,te_d)
def vectorized_result(j):
    """Return a 10-dimensional unit vector with a 1.0 in the jth
    position and zeroes elsewhere. This is used to convert a digit
    (0...9) into a corresponding desired output from the neural
    network."""
    e = numpy.zeros((10, 1))
    e[j-1] = 1.0
    return e
#load_data()

```

```

# -*- coding: utf-8 -*-
"""

```

Created on Thu Apr 9 22:59:21 2015

```

@author: pramod
"""

```

```

import numpy as numpy
import random
import glob
from scipy.misc import *

```

```

def LoadImages(directory, split):
    # get a list of all the picture filenames
    gifs = glob.glob(directory + '/*.gif')
    # uncomment the below line when trying an unknown file
    #extraGif = glob.glob("/media/cosmos/Data/College Notes/M.Tech/Semester 4/Statistical Methods in AI/Project -
Face Recognition/2.gif")
    classMap = {}
    testGIF = []
    allLabels = []
    testLabels = []
    trainingGIF = []
    trainingLabels = []
    for i in range(len(gifs)):
        if random.random() < split:
            trainingGIF.append(gifs[i])
            l = gifs[i].split("/")
            labelName = l[len(l)-1].split(".")[0][-2:]
            trainingLabels.append(labelName)
            allLabels.append(labelName)
        else:
            testGIF.append(gifs[i])
            l = gifs[i].split("/")
            labelName = l[len(l)-1].split(".")[0][-2:]
            testLabels.append(labelName)
            allLabels.append(labelName)
    # uncomment the below 2 lines when trying an unknown file

```

```

#testGIF.append(extraGif[0])
#testLabels.append("un")
#allLabels.append("un")
trainingImgs = numpy.array([imread(i, True).flatten() for i in trainingGIF])
testImgs = numpy.array([imread(i, True).flatten() for i in testGIF])
# creating a list of class labels
allLabels = set(allLabels)
noOfClasses = len(allLabels)
sortedLabels = []
for i in allLabels:
    sortedLabels.append(i)
sortedLabels = sorted(sortedLabels)
# creating a mapping for confusion matrix
j = 0
for i in sortedLabels:
    classMap[i] = j
    j = j + 1
return trainingImgs, testImgs, trainingLabels, testLabels, noOfClasses, classMap
# mean
mu = numpy.mean(data, 0)
# mean adjust the data
ma_data = data - mu
# run SVD
e_faces, sigma, v = numpy.linalg.svd(ma_data.transpose(), full_matrices=False)
# compute weights for each image
weights = numpy.dot(ma_data, e_faces)
return e_faces, weights, mu
def PCA(data):
    # mean
    mu = numpy.mean(data, 0)
    # mean adjust the data
    ma_data = data - mu
    # run SVD
    e_faces, sigma, v = numpy.linalg.svd(ma_data.transpose(), full_matrices=False)
    # compute weights for each image
    weights = numpy.dot(ma_data, e_faces)
    return e_faces, weights, mu
def InputWeight(testData, mu, e_faces):
    ma_data = testData - mu
    weights = numpy.dot(ma_data, e_faces)
    return weights
def load_data():
    inDIR = "/home/cosmos/CSSstuff/SMAI/Project 2 - Classification/input/yalefaces"
    outDIR = "/home/cosmos/CSSstuff/SMAI/Project 2 - Classification/output"
    imgDims = (243, 320)

    split = 0.8
    trainingData, testData, trainingLabels, testLabels, noOfClasses, classMap = LoadImages(inDIR, split)
    e_faces, trainingWeights, mu = PCA(trainingData)
    #print trainingWeights.shape[0]
    #print trainingWeights.shape
    #print mu.shape
    #print e_faces.shape
    '''
    tr_d, va_d, te_d = load_data()
    training_inputs = [np.reshape(x, (784, 1)) for x in tr_d[0]]
    print tr_d[0].shape
    training_results = [vectorized_result(y) for y in tr_d[1]]
    training_data = zip(training_inputs, training_results)
    validation_inputs = [np.reshape(x, (784, 1)) for x in va_d[0]]
    validation_data = zip(validation_inputs, va_d[1])
    test_inputs = [np.reshape(x, (784, 1)) for x in te_d[0]]
    test_data = zip(test_inputs, te_d[1])

```

```

return (training_data, validation_data, test_data)
'''
testWeights=numpy.zeros((len(testData),trainingWeights.shape[1]))
#print testWeights.shape
#print trainingWeights.shape
#print len(testWeights)
for i in range(len(testWeights)):
    #testWeight =InputWeight(testData[i],mu,e_faces)
    testWeights[i]=InputWeight(testData[i],mu,e_faces)

#formatting weights
    training_inputs = [numpy.reshape(x, (len(trainingWeights), 1)) for x in trainingWeights]
    test_inputs= [numpy.reshape(x,(len(trainingWeights) , 1)) for x in testWeights]
#Convert training labels to vectors
trainingLabels=numpy.asarray(trainingLabels)
trainingLabels=trainingLabels.astype(numpy.float)
testLabels=numpy.asarray(testLabels)
testLabels=testLabels.astype(numpy.float)
#decrementing subject labels
for i in range(len(trainingLabels)):
    trainingLabels[i]=trainingLabels[i]-1
for i in range(len(testLabels)):
    testLabels[i]=testLabels[i]-1
#print testLabels
#print vectorized_result(trainingLabels[0])
#print trainingWeights.shape
training_results = [vectorized_result(y) for y in trainingLabels]
#test_results= [vectorized_result(y) for y in testLabels]

tr_d=zip(training_inputs, training_results)
te_d=zip(test_inputs,testLabels)
#print te_d
return (tr_d,te_d)
def vectorized_result(j):
    """Return a 10-dimensional unit vector with a 1.0 in the jth
    position and zeroes elsewhere. This is used to convert a digit
    (0...9) into a corresponding desired output from the neural
    network."""
    e = numpy.zeros((15, 1))
    e[j-1] = 1.0
    return e
#load_data()

```

Neural Network Eigen Utility Code:

```

import eigen_face_loader
trainingData, testData = eigen_face_loader.load_data()
import neuralnetwork
net = neuralnetwork.Network([len(trainingData), 200, 15])
net.SGD(trainingData, 50, 5, .09, 15, test_data = testData)

```

2. Results:

```

Epoch 0: 3 / 27
Epoch 1: 2 / 27
Epoch 2: 2 / 27
Epoch 3: 3 / 27
Epoch 4: 4 / 27
Epoch 5: 4 / 27
Epoch 6: 4 / 27
Epoch 7: 4 / 27
Epoch 8: 3 / 27

```

Epoch 9: 3 / 27
Epoch 10: 3 / 27
Epoch 11: 3 / 27
Epoch 12: 3 / 27
Epoch 13: 3 / 27
Epoch 14: 3 / 27
Epoch 15: 3 / 27
Epoch 16: 4 / 27
Epoch 17: 4 / 27
Epoch 18: 4 / 27
Epoch 19: 4 / 27
Epoch 20: 4 / 27
Epoch 21: 4 / 27
Epoch 22: 4 / 27
Epoch 23: 4 / 27
Epoch 24: 4 / 27
Epoch 25: 4 / 27
Epoch 26: 4 / 27
Epoch 27: 4 / 27
Epoch 28: 4 / 27
Epoch 29: 4 / 27
Epoch 30: 4 / 27
Epoch 31: 4 / 27
Epoch 32: 4 / 27
Epoch 33: 4 / 27
Epoch 34: 4 / 27
Epoch 35: 4 / 27
Epoch 36: 3 / 27
Epoch 37: 3 / 27
Epoch 38: 3 / 27
Epoch 39: 3 / 27
Epoch 40: 3 / 27
Epoch 41: 3 / 27
Epoch 42: 3 / 27
Epoch 43: 3 / 27
Epoch 44: 3 / 27
Epoch 45: 3 / 27
Epoch 46: 3 / 27
Epoch 47: 3 / 27
Epoch 48: 3 / 27
Epoch 49: 3 / 27

Confusion Matrix:

[0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0]
[1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0]
[0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0]
[0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 2, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[1, 0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 0, 0, 1]
[0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1]
[0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
[0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0]

Class 1

Precision : 0.0

Recall : 0.0

Specificity : 93.1034482759

Class 2

Precision : 0.0

Recall : 0.0

Specificity : 100.0

Class 3

Precision : 0.0

Recall : 0.0

Specificity : 96.4285714286

Class 4

Precision : 25.0

Recall : 33.3333333333

Specificity : 89.6551724138

Class 5

Precision : 0.0

Recall : 0.0

Specificity : 96.4285714286

Class 6

Precision : 0.0

Recall : 0.0

Specificity : 93.1034482759

Class 7

Precision : 33.3333333333

Recall : 33.3333333333

Specificity : 92.8571428571

Class 8

Precision : 0.0

Recall : 0.0

Specificity : 96.4285714286

Class 9

Precision : 50.0

Recall : 50.0

Specificity : 96.2962962963

Class 10

Precision : 0.0

Recall : 0.0

Specificity : 93.1034482759

Class 11

Precision : 0.0

Recall : 0.0

Specificity : 93.1034482759

Class 12

Precision : 0.0

Recall : 0.0

Specificity : 90.0

Class 13

Precision : 0.0

Recall : 0.0

Specificity : 96.4285714286

Class 14

Precision : 0.0

Recall : 0.0

Specificity : 96.4285714286

Class 15

Precision : 0.0

Recall : 0.0

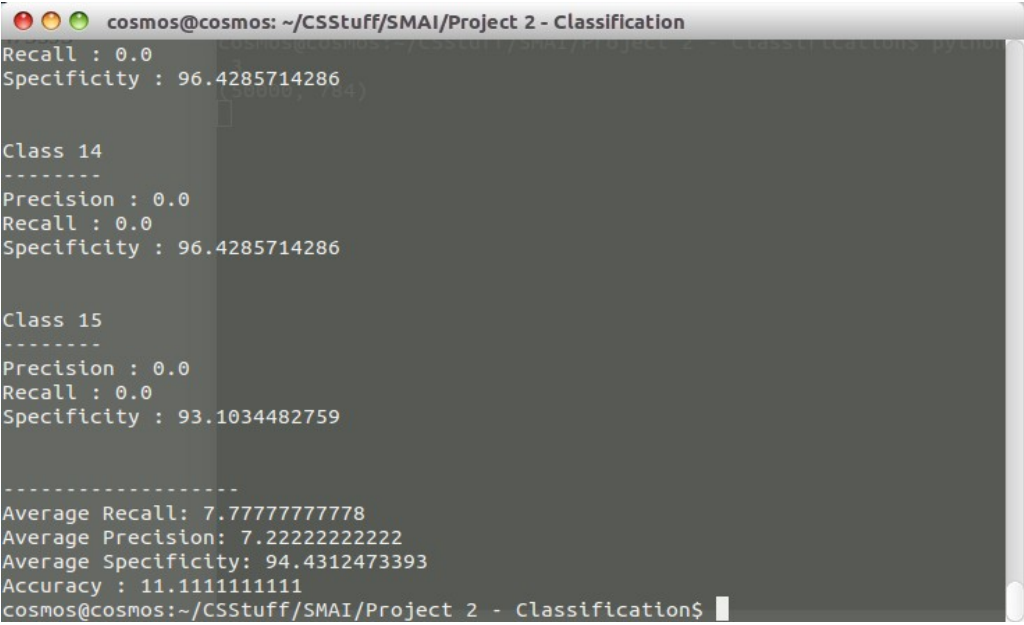
Specificity : 93.1034482759

Average Recall: 7.7777777778

Average Precision: 7.2222222222

Average Specificity: 94.4312473393

Accuracy : 11.1111111111



```
cosmos@cosmos: ~/CSStuff/SMAI/Project 2 - Classification
Recall : 0.0
Specificity : 96.4285714286

Class 14
-----
Precision : 0.0
Recall : 0.0
Specificity : 96.4285714286

Class 15
-----
Precision : 0.0
Recall : 0.0
Specificity : 93.1034482759

-----
Average Recall: 7.7777777778
Average Precision: 7.2222222222
Average Specificity: 94.4312473393
Accuracy : 11.1111111111
cosmos@cosmos:~/CSStuff/SMAI/Project 2 - Classification$
```

For MNIST Dataset:

1. Souce Code

mnist_loader :

```
# -*- coding: utf-8 -*-  
"""
```

Created on Thu Apr 9 21:09:06 2015

```
@author: pramod  
"""
```

```
"""
```

```
mnist_loader
```

```
~~~~~
```

A library to load the MNIST image data. For details of the data structures that are returned, see the doc strings for ``load_data`` and ``load_data_wrapper``. In practice, ``load_data_wrapper`` is the function usually called by our neural network code.

```
"""
```

```
#### Libraries
```

```
# Standard library
```

```
import cPickle
```

```
import gzip
```

```
# Third-party libraries
```

```
import numpy as np
```

```
def load_data():
```

```
    """Return the MNIST data as a tuple containing the training data,  
    the validation data, and the test data.
```

The ``training_data`` is returned as a tuple with two entries.

The first entry contains the actual training images. This is a numpy ndarray with 50,000 entries. Each entry is, in turn, a numpy ndarray with 784 values, representing the $28 * 28 = 784$ pixels in a single MNIST image.

The second entry in the ``training_data`` tuple is a numpy ndarray containing 50,000 entries. Those entries are just the digit values (0...9) for the corresponding images contained in the first entry of the tuple.

The ``validation_data`` and ``test_data`` are similar, except each contains only 10,000 images.

This is a nice data format, but for use in neural networks it's helpful to modify the format of the ``training_data`` a little. That's done in the wrapper function ``load_data_wrapper()`` , see below.

```
"""
```

```
f = gzip.open('/home/cosmos/CSStuff/SMAI/Project 2 - Classification/input/mnist.pkl.gz', 'rb')  
training_data, validation_data, test_data = cPickle.load(f)  
f.close()  
return (training_data, validation_data, test_data)
```

```
def load_data_wrapper():
```

```
    """Return a tuple containing ``(training_data, validation_data,  
    test_data)``. Based on ``load_data``, but the format is more  
    convenient for use in our implementation of neural networks.
```

In particular, ``training_data`` is a list containing 50,000 2-tuples ``(x, y)``. ``x`` is a 784-dimensional `numpy.ndarray` containing the input image. ``y`` is a 10-dimensional `numpy.ndarray` representing the unit vector corresponding to the correct digit for ``x``.

``validation_data`` and ``test_data`` are lists containing 10,000 2-tuples ``(x, y)``. In each case, ``x`` is a 784-dimensional `numpy.ndarray` containing the input image, and ``y`` is the corresponding classification, i.e., the digit values (integers) corresponding to ``x``.

Obviously, this means we're using slightly different formats for the training data and the validation / test data. These formats turn out to be the most convenient for use in our neural network code. """

```
tr_d, va_d, te_d = load_data()
training_inputs = [np.reshape(x, (784, 1)) for x in tr_d[0]]

training_results = [vectorized_result(y) for y in tr_d[1]]
training_data = zip(training_inputs, training_results)
validation_inputs = [np.reshape(x, (784, 1)) for x in va_d[0]]
validation_data = zip(validation_inputs, va_d[1])
test_inputs = [np.reshape(x, (784, 1)) for x in te_d[0]]
test_data = zip(test_inputs, te_d[1])
return (training_data, test_data)
```

```
def vectorized_result(j):
    """Return a 10-dimensional unit vector with a 1.0 in the jth
    position and zeroes elsewhere. This is used to convert a digit
    (0...9) into a corresponding desired output from the neural
    network. """
    e = np.zeros((10, 1))
    e[j] = 1.0
    return e
```

MNIST Utility Function Source Code:

```
import mnist_loader
trainingData, test_data = mnist_loader.load_data_wrapper()
import copyCode
net = copyCode.Network([784, 30, 10])
net.SGD(trainingData, 30, 10, 3.0, 10, test_data)
```

2. Results

```
Epoch 0: 9104 / 10000
Epoch 1: 9211 / 10000
Epoch 2: 9262 / 10000
Epoch 3: 9321 / 10000
Epoch 4: 9358 / 10000
Epoch 5: 9370 / 10000
Epoch 6: 9375 / 10000
Epoch 7: 9395 / 10000
Epoch 8: 9390 / 10000
Epoch 9: 9399 / 10000
```


Confusion Matrix:

[932, 0, 5, 0, 1, 4, 6, 1, 1, 8]
[0, 1108, 0, 1, 2, 3, 1, 5, 1, 4]
[4, 3, 959, 12, 2, 2, 2, 27, 7, 1]
[5, 2, 9, 939, 2, 23, 4, 10, 12, 15]
[2, 0, 7, 1, 927, 3, 7, 9, 4, 27]
[7, 1, 5, 22, 2, 822, 9, 1, 9, 8]
[11, 7, 12, 3, 12, 7, 913, 1, 3, 0]
[1, 1, 12, 7, 2, 1, 0, 956, 7, 15]
[13, 13, 19, 19, 8, 25, 16, 6, 928, 16]
[5, 0, 4, 6, 24, 2, 0, 12, 2, 915]

Class 1

Precision : 95.1020408163

Recall : 97.2860125261

Specificity : 99.473453269

Class 2

Precision : 97.6211453744

Recall : 98.4888888889

Specificity : 99.6972754793

Class 3

Precision : 92.9263565891

Recall : 94.1118743867

Specificity : 99.1990344525

Class 4

Precision : 92.9702970297

Recall : 91.9686581783

Specificity : 99.2225142357

Class 5

Precision : 94.399185336

Recall : 93.9209726444

Specificity : 99.3974583699

Class 6

Precision : 92.1524663677

Recall : 92.776523702

Specificity : 99.2430795848

Class 7

Precision : 95.3027139875

Recall : 94.2208462332

Specificity : 99.5072273325

Class 8

Precision : 92.9961089494

Recall : 95.4091816367
Specificity : 99.2101799035

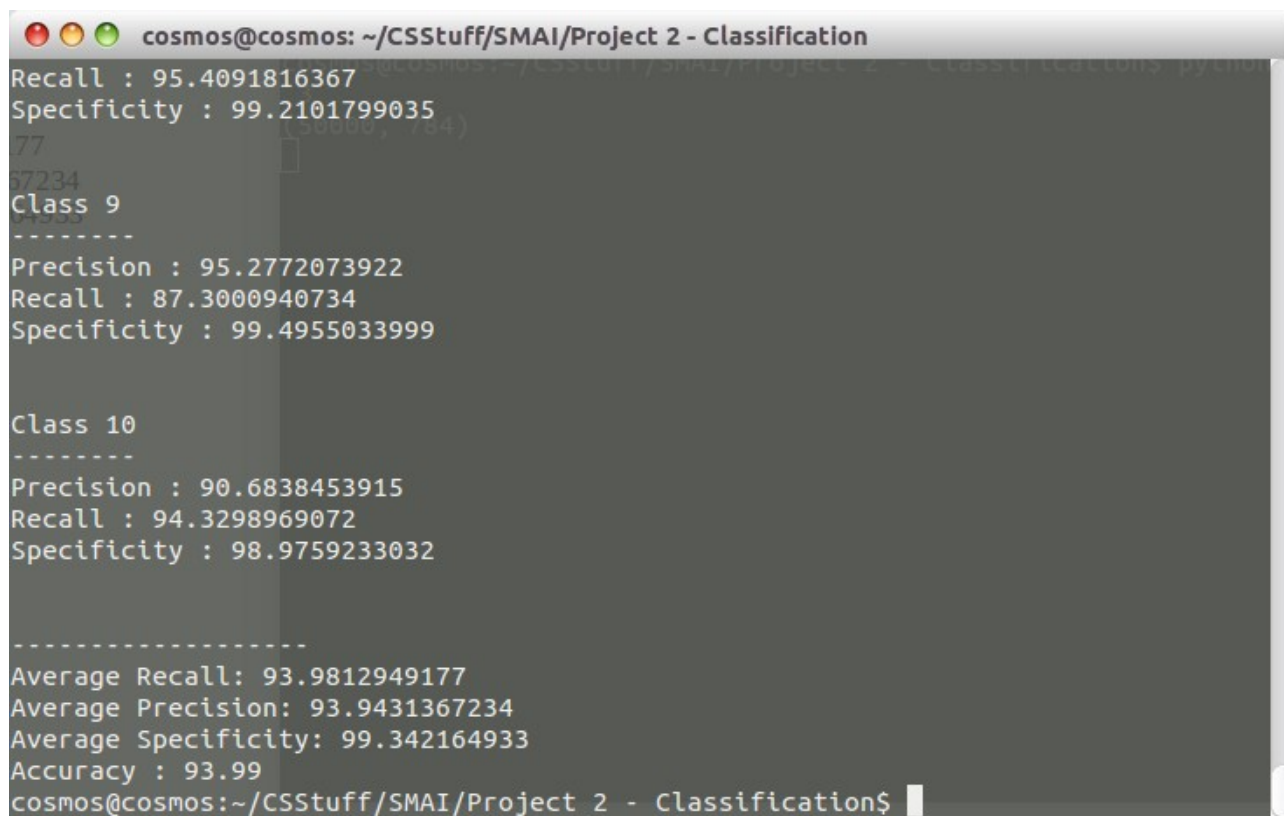
Class 9

Precision : 95.2772073922
Recall : 87.3000940734
Specificity : 99.4955033999

Class 10

Precision : 90.6838453915
Recall : 94.3298969072
Specificity : 98.9759233032

Average Recall: 93.9812949177
Average Precision: 93.9431367234
Average Specificity: 99.342164933
Accuracy : 93.99

A terminal window titled "cosmos@cosmos: ~/CSStuff/SMAI/Project 2 - Classification" displays the same classification metrics as the previous blocks. The output is as follows:

```
cosmos@cosmos: ~/CSStuff/SMAI/Project 2 - Classification
Recall : 95.4091816367
Specificity : 99.2101799035
-----
Class 9
-----
Precision : 95.2772073922
Recall : 87.3000940734
Specificity : 99.4955033999

Class 10
-----
Precision : 90.6838453915
Recall : 94.3298969072
Specificity : 98.9759233032

-----
Average Recall: 93.9812949177
Average Precision: 93.9431367234
Average Specificity: 99.342164933
Accuracy : 93.99
cosmos@cosmos:~/CSStuff/SMAI/Project 2 - Classification$
```

Comparisons Between Different Methods:

| Dataset / Methods | K- Nearest Neighbours | SVM | Multilayer Feedforwards Neural Network |
|-------------------|-----------------------|-----------------|--|
| Yale Dataset | 86.1221321373 % | 90.0948621554 % | 11.11 % |
| MNIST Dataset | 82.15 % | 94.35 % | 93.99% |